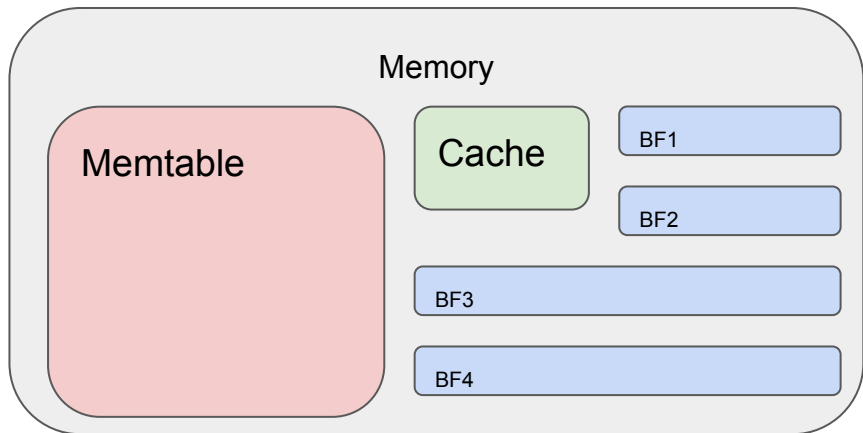





Optimal Memory Allocation

Mali Akmanalp, Sophie Hilgard, Andrew Ross

Motivation



What does it depend on



What does it depend on



We know the right size matters:

Monkey: Optimal Navigable Key-Value Store

Niv Dayan
Harvard University


dayan@seas.harvard.edu

Manos Athanassoulis
Harvard University

manos@seas.harvard.edu

Stratos Idreos
Harvard University


stratos@seas.harvard.edu




But can we consider the LSM tree as a whole?



But can we consider the LSM tree as a whole?



But can we consider the LSM tree as a whole?



But can we consider the LSM tree as a whole?



But can we consider the LSM tree as a whole?



But can we consider the LSM tree as a whole?

(+ the query distribution?)



How much do optimal sizes depend on the query distribution?

How well do we need to know the query distribution?



High-level strategy



Modeling

+

Simulation

+

Benchmarks



Make sure they all agree


High-level strategy

Then:

Learn what matters



Annealing + Inference?



Implement
in RocksDB









Preliminary Results

Simulation

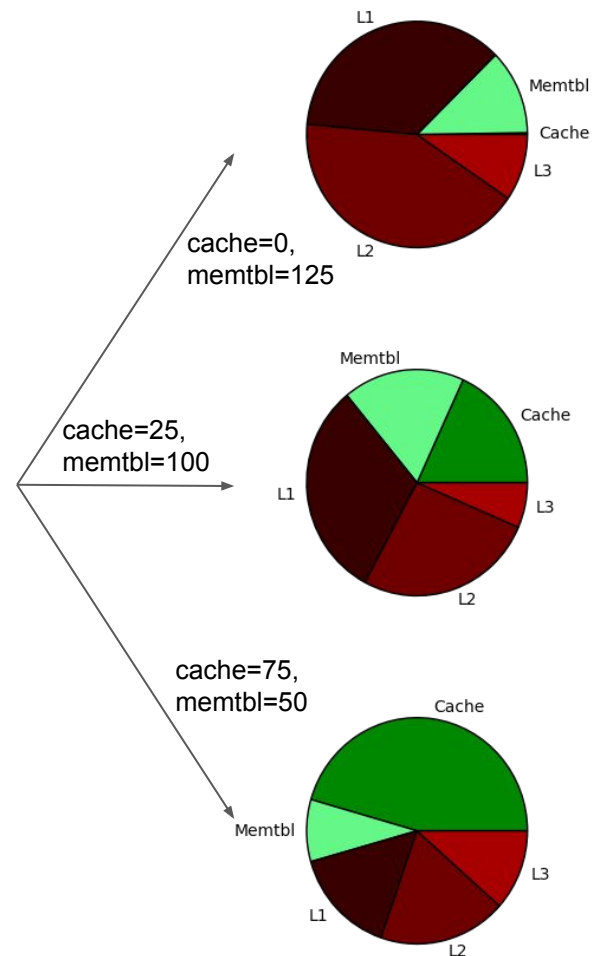

Simulated tree with cache, memtable, layers, and bloom filters w/ configurable sizes.

Same LSM tree architecture
+ Different query distributions
= Very different outcomes



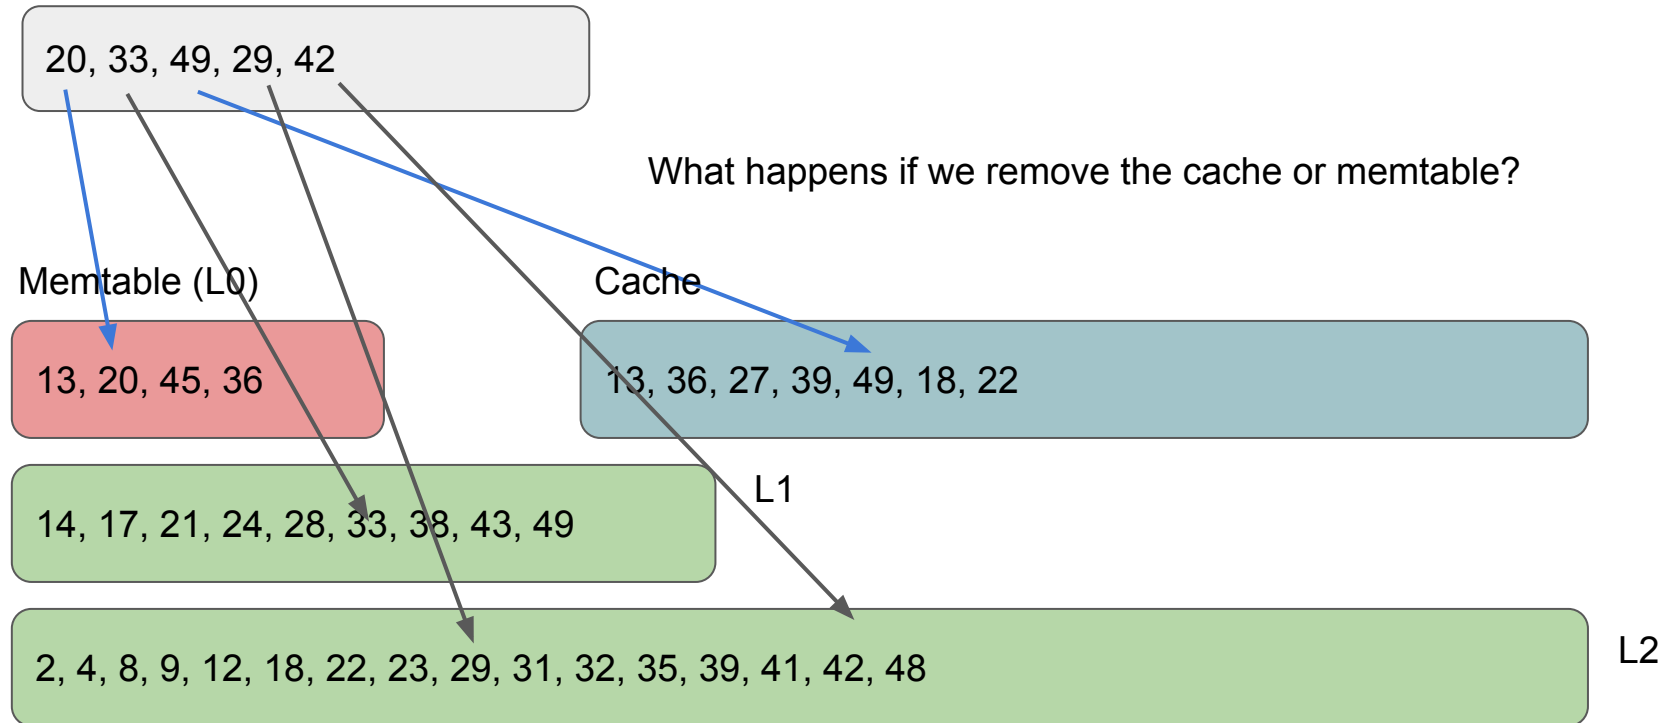
Simulation

Different LSM tree architectures are best for different query distributions.



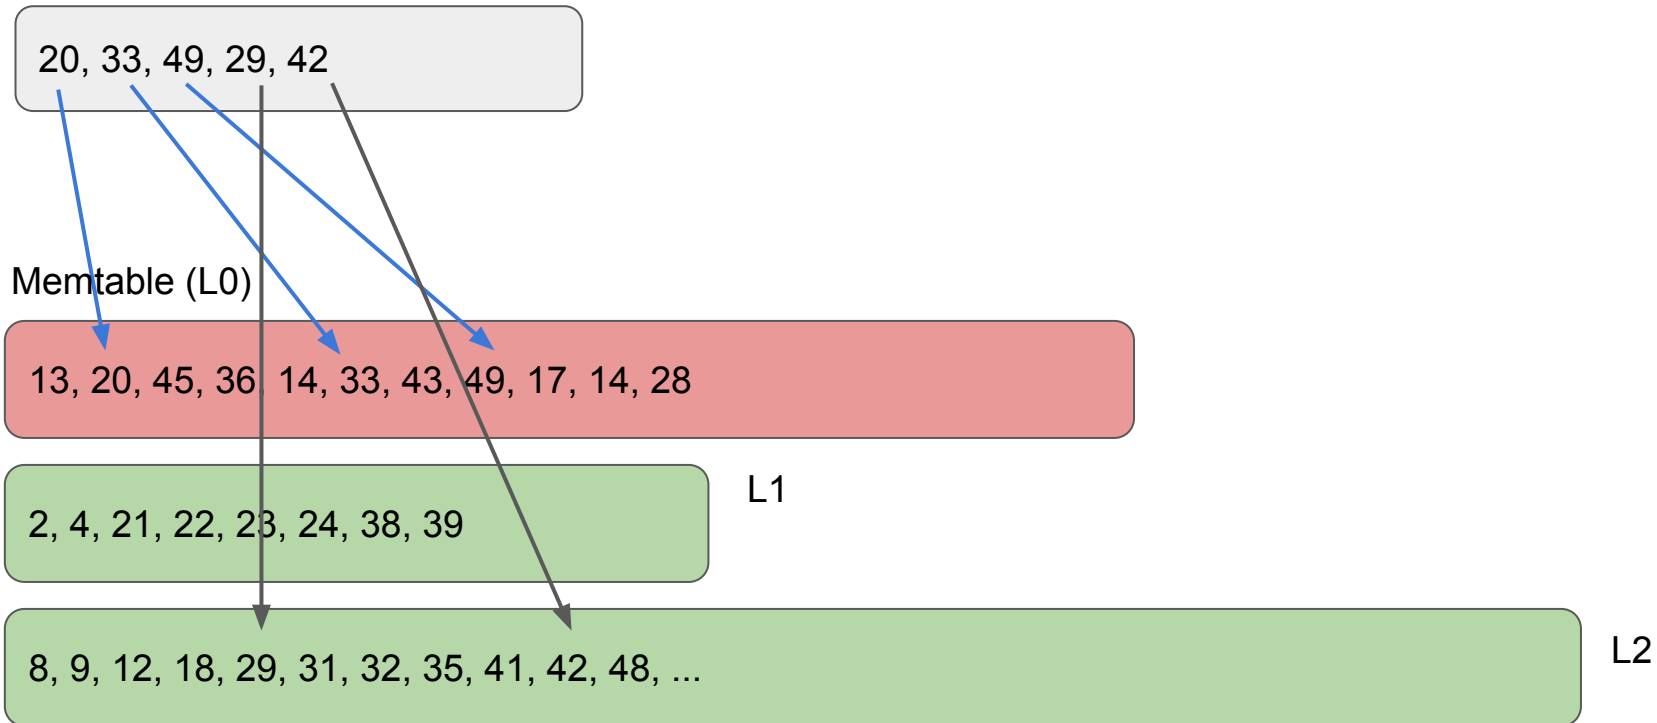
Modeling

Randomly Distributed Requests



Modeling

Randomly Distributed Requests



Modeling

Randomly Distributed Requests

20, 33, 49, 29, 42

Cache

13, 36, 27, 39, 49, 18, 22, 33, 48, 14, 24, 32

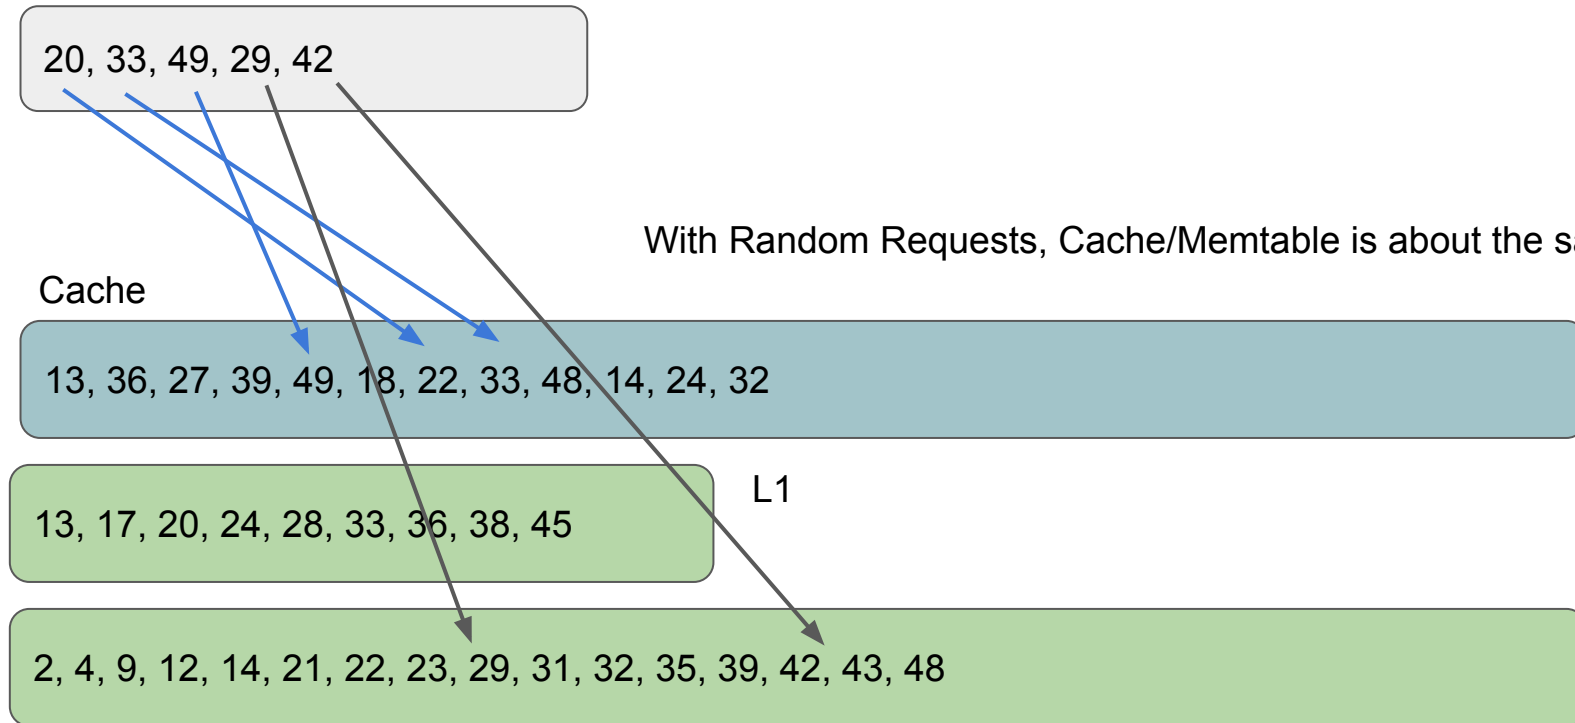
13, 17, 20, 24, 28, 33, 36, 38, 45

L1

2, 4, 9, 12, 14, 21, 22, 23, 29, 31, 32, 35, 39, 42, 43, 48

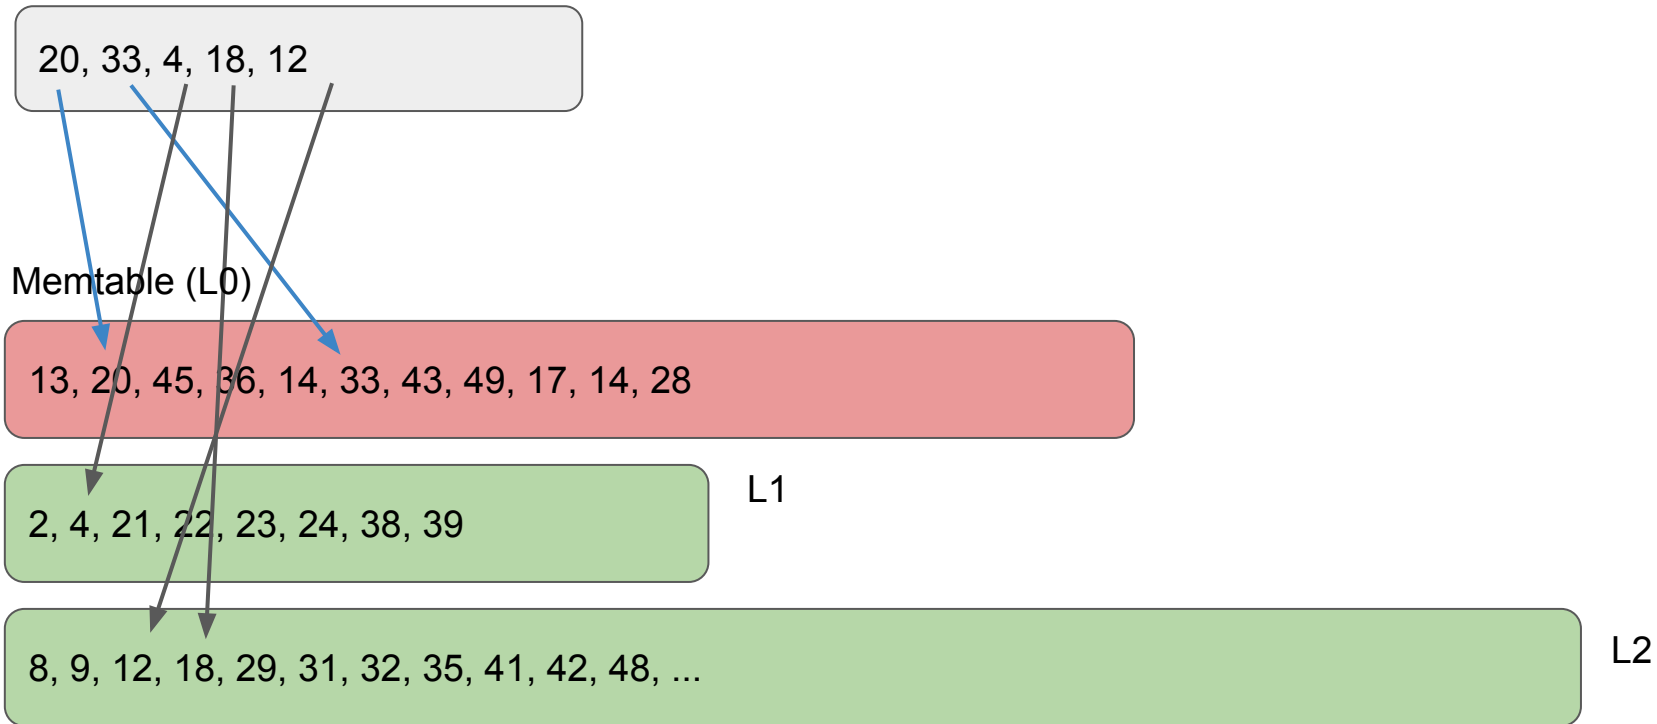
L2

With Random Requests, Cache/Memtable is about the same



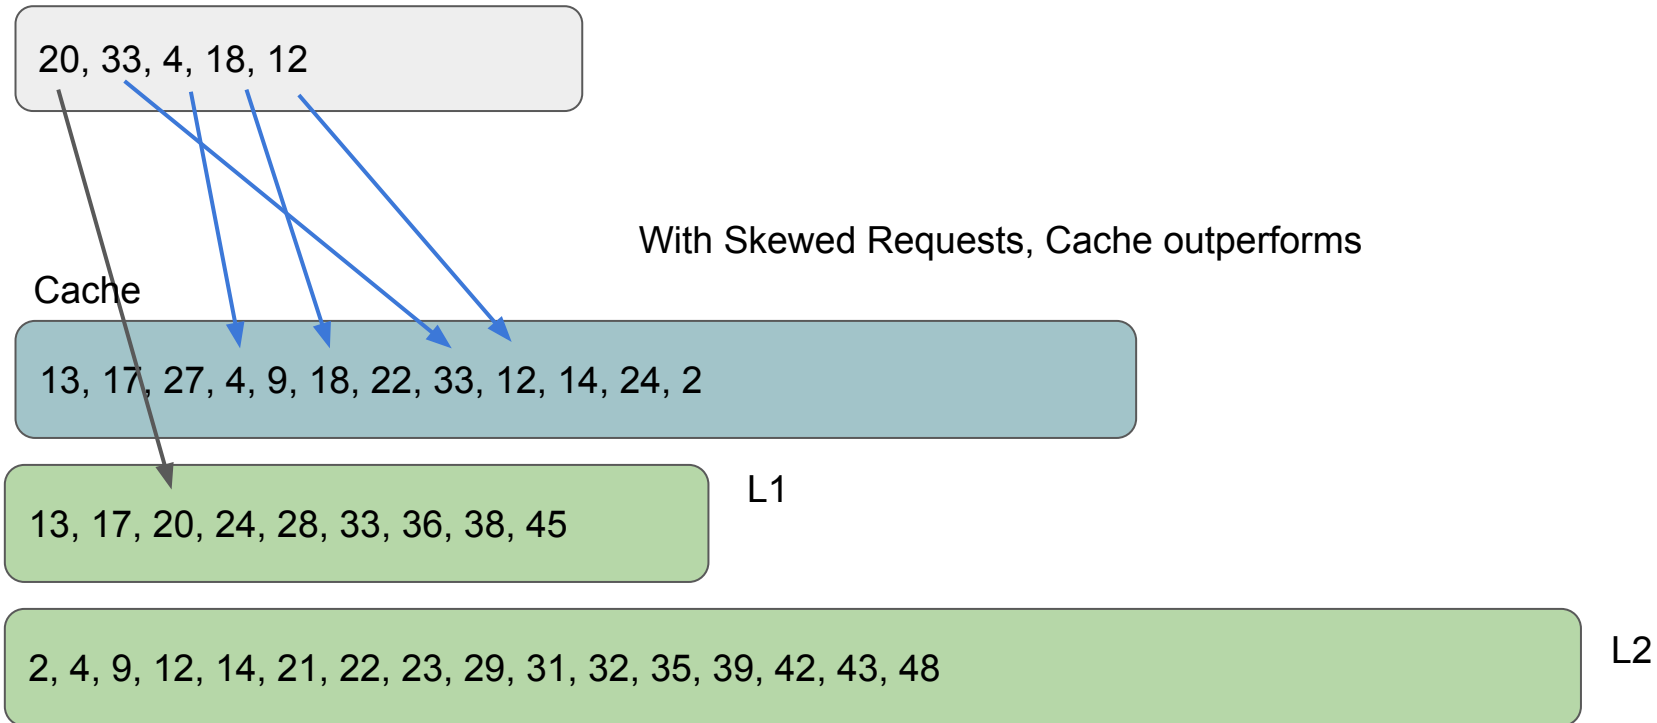
Modeling

Skewed Requests (More Heavily Weighted 1-20)




Modeling

Skewed Requests (More Heavily Weighted 1-20)

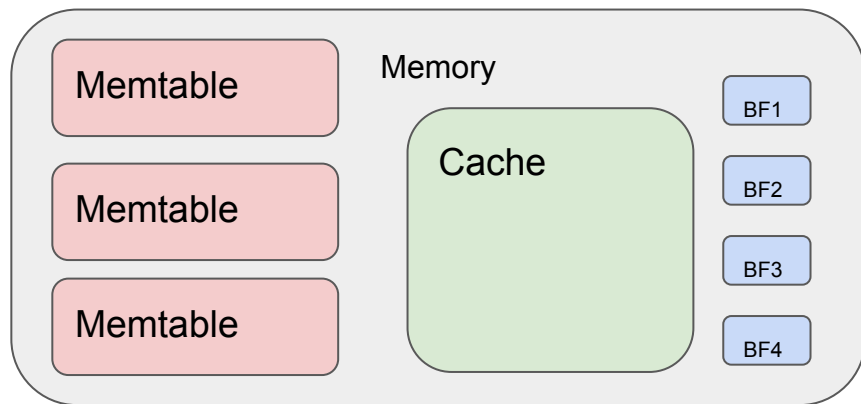


Parameter space for experiment



- Memtable size
- Cache size
- Bloom filter size for each layer

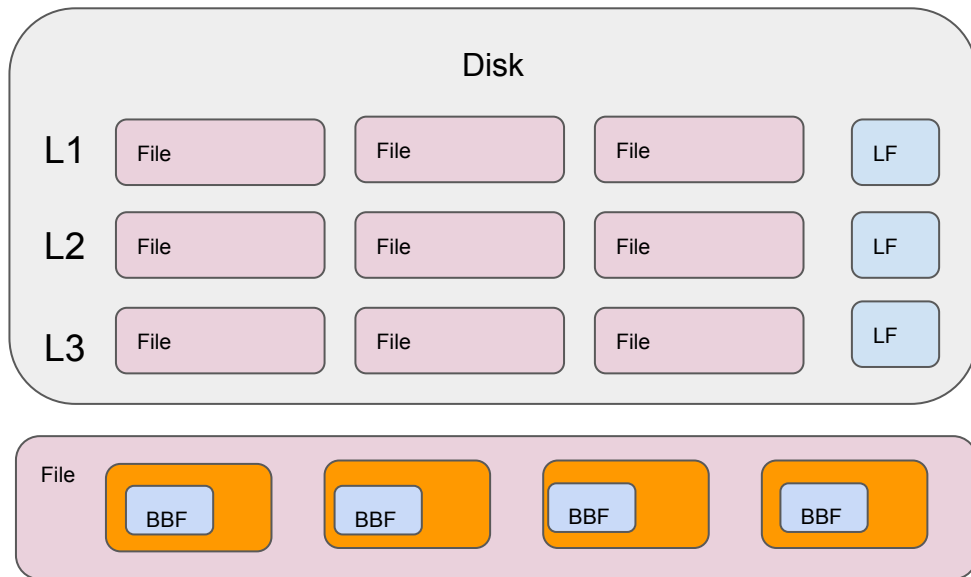
Memtables in RocksDB: “write buffers”



- **write_buffer_size:**
memtable size
- **max_write_buffer_number:**
number of memtables
- **min_write_buffer_number_to_merge**
configure how often
memtables get written

For us, memtable size is the total of all the write buffers.
We probably want the Vector memtable.

Bloom Filters in RocksDB



- **Block based bloom filter**
For each block
- **Layer bloom filter:**
For a full layer

For us, bloom filter size is the total of all layer filters.

Caches in RocksDB




For us, cache size is the total block cache.

- **No key cache?**
- **Persistent read cache:**
On-disk cache to look up blocks and files
- **Block cache:**
Uncompressed blocks from files: LRU / Time
- **OS Page cache?**
Unsure if we should try to bypass or include

Measurements

- Simulation Cache:
 - Neat feature that allows you to simulate cache behavior with different sizes on a real workload
- CreateDBStatistics()
 - Compaction Stats: Read/Write/Moved (GB)
 - General Stats: number / time spent of writes?
 - Block cache:
 - Hit/Miss rate
 - Layers:
 - Hit/miss for memtable
 - Hit for L1/L2/L2+ (interesting that it's not for all layers)
 - Bloom Filter:
 - Bloom filter "usefulness"

Problems and Bottlenecks



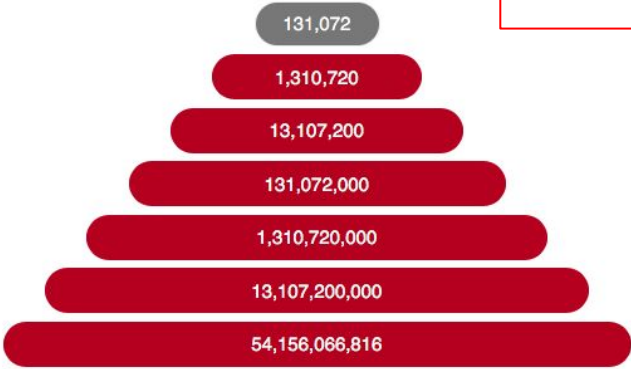
MONKEY focused on
changing just one thing.

How do we define success?

MONKEY VS. STATE-OF-THE-ART

		FALSE POSITIVE RATES	
		State-Of-The-Art	Monkey
Buffer:	Level		
	0	--	--
	1	0.82%	0.000028%
	2	0.82%	0.00028%
	3	0.82%	0.0028%
	4	0.82%	0.028%
	5	0.82%	0.28%
	6	0.82%	1.2%

ENTRIES PER LEVEL



Can we show similar results for different workloads and more memory options?

Monkey lookups are 3.32x faster!

Lookup cost:	0.049 I/Os	0.015 I/Os
Update cost:	0.105 I/Os	0.105 I/Os
Main memory:	80.002 GB	80.002 GB

Preliminary Conclusions

Optimization is hard

How to draw an owl

1.



1. Draw some circles

2.



2. Draw the rest of the finished owl

But important



RocksDB offers good building blocks



but it might be hard to balance.

Actual quote: “Even we as RocksDB developers don't fully understand the effect of each configuration change”

Math + simulation + experimentation seems tractable.

Math + simulation + experimentation seems tractable.

hopefully.